

# Improving RAG Applications

For MIT Generative AI Course

---

Jason Liu

February 25, 2025

Independent Consultant

# Table of Contents

About Me

What are we doing here?

Practical Implementation Strategies

- Evaluation Methods

- Query Segmentation

- Specialized Retrieval

- Feedback Collection

- Implementation Roadmap

## About Me

---

# About Me

*Goal: To showcase my diverse experience – feel free to ask questions about any area during Q&A*

- Independent Consultant & Staff-level ML Engineer & Educator
  - Meta, Stitchfix, NYU from (2016-2023)
- University of Waterloo
  - B.Math in Mathematical Physics & Computational Mathematics
  - Minor in Statistics (Class of 2017)
- Creator of [Instructor](#) - Python library for structured LLM outputs
  - 9500+ GitHub stars
  - 1.5M+ monthly downloads
  - Cited by OpenAI as inspiration for structured output feature
  - Popular 'Pydantic is all you need' Talk from AI Conference
- a16z Scout & [Angel Investor](#) & Startup Advisor

- Transitioned to independent consulting due to RSI (2022)
- Do I want to get 20% of my coding back, or find 100x more leverage?
- Now focused on:
  - Teaching teams to work with AI and be quantitative
  - Writing more 'popular ai' content
  - Advisory work for early stage startups
  - Open source projects, independent research

## Work & Consulting & Advisory Engagements

Client	Contact	Industry
Zapier	VP of Product	Automation
HubSpot	GM	Sales & Marketing
Enterpret	CTO	Analytics
Tensorlake	CEO	Data
Limitless AI	CTO	AI
Trunk Tools	VP Eng	Construction
Naro	CTO	Sales & Marketing

Additionally, I've worked with innovative startups including [New Computer](#), [Sandbar](#), [Dunbar](#), [Bytebot](#), [Kay.ai](#), [Raycast](#), [Weights & Biases](#), [Modal Labs](#), [Timescale](#), and [Pydantic](#) on various technical and strategic initiatives.

# Where My Students Come From

Company	Industry
OpenAI	AI Research & Development
Anthropic	AI Research & Development
Google	Search Engine
Salesforce	Customer Relationship Management Software
Microsoft	Software, Cloud Computing
Amazon	E-commerce, Cloud Computing
Zapier	Automation Software
Adobe	Software, Creative Tools
Accenture	Consulting, Technology Services
McKinsey & Company	Management Consulting
Bain & Company	Consulting
PwC	Professional Services
Cisco	Networking Technology
Electronic Arts	Gaming
Shopify	E-commerce Platform

**What are we doing here?**

---



## ■ Learning Objectives

- Developing durable AI knowledge that outlasts specific technical implementations
- Understanding ML systems as continuously evolving products rather than "deploy once and forget"
- Recognizing the parallels between recommendation systems and retrieval systems
- Identifying valuable business applications through effective data analysis

# Setting the Context

## ■ Why This Matters Now

- Democratization of AI tools means the competitive advantage comes from thinking deeply
- Growing gap between research capabilities and business implementation
- Science now drives product development, reversing traditional patterns
- Opportunity for individual contributors to have outsized impact through thoughtful implementation

## ■ Interactive Format

- This group is quite diverse, so I'll try to keep it broad
- The goal is to seed you with good questions, rather than dump information
- We'll leave plenty of time for questions about AI, Business, and Career paths

## Key Questions to Consider

- How has machine learning research and implementation evolved from 2015 to 2025?
- What behavioral practices should teams adopt when working with AI systems?
- How do we identify economically valuable AI applications?
- What's the right balance between research exploration and product implementation?
- How do we design systems that can evolve effectively over time?
- When should you join established labs versus work independently?
- How can individuals and small teams achieve leverage without large resources?
- What skills matter most in the AI era? (Hint: thinking & coding)

# Three Key Arcs We'll Explore

- Technical: From Recommendation to Retrieval Systems
  - The surprising similarities in architecture and challenges
  - Why understanding these parallels helps build better systems
- Organizational: Effective AI Implementation
  - The importance of observability and measurement
  - Balancing unified systems vs. specialized subsystems
- Personal: Career Considerations in AI
  - Information synthesis as a durable skill
  - How AI changes team dynamics and individual contributions

# Practical Implementation Strategies

---

## Synthetic Data Generation (1/2)

- Use LLMs to generate domain-specific questions from your documents
- Create diverse query types:
  - Factual: "What was the revenue in Q2 2023?"
  - Comparative: "How did Q2 performance compare to Q1?"
  - Analytical: "What factors contributed to the margin decline?"

**Benefits:** Controllable test data, covers edge cases, identifies blind spots

## Synthetic Data Generation (2/2)

- Include edge cases and known failure modes
- Example prompt:

"Generate 10 questions a financial analyst might ask about this earnings report, including questions about revenue trends, profitability metrics, and forward guidance."

- Aim for 100-200 diverse test examples across categories

# Retrieval Metrics (1/2)

- **Recall@k**: Percentage of relevant documents in top k results
  - Critical for RAG - can't generate from missing information
  - Target 80-90% recall before focusing on generation quality
- **MRR (Mean Reciprocal Rank)**: Position of first relevant document
  - Higher weight to documents appearing earlier in results
  - Useful for prioritizing most relevant content



## Retrieval Metrics (2/2)

- **Latency:** Response time for retrieval operations
  - Critical for user experience and scaling
  - Balance between accuracy and speed
- **Coverage:** Percentage of query types that can be answered
  - Identifies gaps in retrieval capabilities
  - Helps prioritize new index or tool development

**Key insight:** Focus on recall first!  
Poor recall = ceiling on overall quality

### Evaluation Systems: The Foundation for Improvement

- **Start with synthetic data** to enable rapid testing cycles
- **Focus on retrieval quality first** before optimizing generation
- **Build fast evaluation pipelines** that run in under a minute
- **Collect 100-200 diverse test examples** across query categories
- **Prioritize recall** as the leading indicator of system quality
- **Log and analyze failures** systematically to identify patterns

**Remember:** You can't generate accurate answers from missing information

## Test Size

- 100-200 test examples across categories
- 20-30 examples per major query type
- Include examples from each domain area

## Failure Analysis

- Log all failures for systematic analysis
- Store full context of failed retrievals
- Look for patterns in failure modes

## Automation

- Automate evaluation to run in under 1 minute
- Fast feedback loops enable rapid iteration
- Run after every significant change

# Case Study: Due Diligence Summaries & Missing Experts

## The Challenge

- Consulting firm conducting M&A due diligence
- AI summaries missed half of relevant expert quotes
- Only 3 of 6 experts cited on key topics (50% recall)
- Low recall undermined consultant confidence

## The Solution

- Created "ground truth" set tagging who said what
- Refined chunk boundaries and indexing strategy
- Recall improved from 50% to 90%
- Systematic evaluation led to rapid improvement

**Key Takeaway:** Small "gold set" evaluation + refined chunking dramatically improves system credibility

# Segmentation: The Foundation for Improvement

**Why segment queries?** To identify specific areas for improvement

- Moving beyond basic metrics (recall/precision)
- Identifying patterns in user behavior
- Prioritizing development efforts based on data
- Tracking performance across different query types

**Key insight:** Summary statistics often mask important patterns

# Query Segmentation Approaches (1/2)

- **Intent-based:** What is the user trying to accomplish?
  - Information seeking vs. task completion
  - Exploratory vs. targeted queries
  - Example: "Tell me about X" vs. "How do I do Y?"
- **Domain-based:** Which knowledge area does this touch?
  - Subject matter categories
  - Technical vs. business vs. compliance
  - Example: Financial metrics vs. operational details

## Query Segmentation Approaches (2/2)

- **Complexity-based:** Simple lookup vs. multi-step reasoning
  - Single fact retrieval vs. synthesis across documents
  - Explicit vs. implicit information
  - Example: Direct value lookup vs. trend analysis
- **Data-type:** Text, table, image, code, or multi-modal
  - Different content formats require different handling
  - Example: Narrative text vs. tabular financial data

# Categorizing Issues: Inventory vs. Capabilities

## Inventory Issues

- Missing content in knowledge base
- Incomplete data sources
- Outdated information
- Gaps in document coverage

## Capability Issues

- System's functional limitations
- Missing metadata extraction
- Lack of structured filtering
- Insufficient query understanding

**Different solutions:** Inventory → expand corpus  
Capabilities → build new functionality



# Detecting Inventory vs. Capability Gaps

## Inventory Gap Signals

- Low cosine similarities
- No results from lexical search
- LLM refusing to answer
- Returned chunks never cited
- Broken data pipelines

## Capability Gap Examples

- Time-based filtering needs
- Comparison across documents
- Structured data extraction
- Missing metadata fields
- Need for specialized indices

**Example:** "Latest contract modifications" requires both inventory (recent docs) and capability (time filtering)

# Implementing Query Classification

- Use a classifier prompt to categorize each query
  - Chain-of-thought prompting improves accuracy
  - Allow multiple categories per query when relevant

## Prompt Example:

"Analyze this query and determine which category it belongs to. Explain your reasoning before giving your final answer."

# Tracking Segmentation Performance

- Track performance metrics per segment
  - Separate dashboards for each major category
  - Compare performance across segments
- Identify segments with highest volume  $\times$  lowest performance
  - Focus improvements on high-impact areas
  - Prioritize based on business value

# Prioritization Framework

## Prioritization Formula

- Impact of answering this type of question
- Query volume for the segment
- Likelihood of success (can we solve it?)

$$\text{Priority} = \text{Impact} \times \text{Volume} \times P(\text{Success})$$

## Decision Matrix

- High satisfaction + High volume = Maintain
- High satisfaction + Low volume = Promote
- Low satisfaction + Low volume = Phase out
- Low satisfaction + High volume = Focus here!

### Query Segmentation: The Path to Targeted Improvements

- **Segment queries** by intent, domain, complexity, and data type
- **Distinguish between inventory and capability gaps**
  - Inventory: Missing content → Add more data
  - Capability: System limitations → Build new features
- **Track performance by segment** to identify specific weaknesses
- **Prioritize improvements** based on impact, volume, and feasibility
- **Focus on high-volume, low-satisfaction segments** first

**Remember:** Summary statistics often mask important patterns

# Financial Query Types (1/2)

## ■ Numerical extraction

- Finding specific values in financial statements
- Example: "What was the EPS in Q2 2023?"
- Requires precise table extraction and entity recognition

## ■ Trend analysis

- Identifying patterns over time in financial data
- Example: "How has the gross margin changed over the last 4 quarters?"
- Requires time-series data and comparative analysis

# Financial Query Types (2/2)

## ■ Comparative analysis

- Comparing entities, periods, or metrics
- Example: "How does Company X's ROI compare to industry average?"
- Requires multi-document retrieval and normalization

## ■ Risk assessment

- Evaluating potential issues or concerns
- Example: "What are the key risk factors mentioned in the report?"
- Requires understanding of risk terminology and context

# Real-World Example: Construction Project

## Initial Analysis

- 80
- 20
- Document search had high satisfaction
- Schedule queries had low satisfaction

## Time-Based Analysis

- New users started with schedule queries
- Poor results led to behavior change
- Users learned to use document search as workaround
- Masked the actual problem

**Solution:** Built specialized data extraction for dates and schedules  
Highlighted new capability →  
user behavior changed back



# Case Study: Scheduling & Learned User Behavior

## The Challenge

- Construction management platform
- New users struggled with scheduling queries
- Veteran users learned workarounds
- Only 20% success rate for new users on scheduling queries

## The Solution

- Split "document search" vs. "scheduling" segments
- Created specialized scheduling index
- Extracted due dates and milestones
- Announced new feature to users

**Key Takeaway:** Segmentation + dedicated meta-data indices dramatically improve user satisfaction

# Hybrid Search: Combining Approaches

## Lexical Search

- Exact keyword matching
- Based on BM25, TF-IDF
- Great for precise terms
- Example: Elasticsearch

## Semantic Search

- Meaning-based matching
- Uses embeddings
- Great for concepts
- Handles synonyms

**Hybrid Search:** Combines  
strengths of both approaches

# Hybrid Search: Implementation

- Weight results based on query characteristics
  - Adjust weights dynamically based on query type
  - Use query classifier to determine optimal weights

## When to favor each approach:

- **Lexical:** Specific terms, codes, IDs, exact phrases
- **Semantic:** Concepts, themes, topics, intentions

**Example:** "Q2 2023 revenue" → 70  
"Growth strategy reasons" → 20

# Metadata Enhancement: Extraction

- Extract structured data during ingestion

## Temporal

- Dates
- Time periods
- Fiscal quarters
- Years

## Entities

- Companies
- People
- Products
- Categories

## Document

- Doc type
- Sections
- Importance
- Source

**Key insight:** Rich meta-data enables powerful filtering

# Metadata Enhancement: Filtering

- Create filters based on document attributes
  - Filter by date range, document type, entity
  - Combine filters with semantic search
  - Improve precision without hurting recall
- Enable faceted search capabilities
  - Allow users to narrow results by metadata
  - Provide context-aware filtering options

**Example:** "Revenue for Q2  
2023 in North America region"  
`embedding_search(query) AND  
date="Q2 2023" AND region="NA"`

# Specialized Indices: Content Types

- Create separate indices for different content types

## Text Documents

- Reports
- Articles
- Narratives
- Analysis

## Tabular Data

- Financial statements
- Metrics tables
- KPI dashboards

## Visual Content

- Charts
- Graphs
- Diagrams

## Structured Data

- SQL
- JSON
- XML
- Code

# Specialized Indices: Chunking Strategies

- Optimize chunking strategy per content type

## Text

- Semantic paragraphs
- Fixed-size chunks
- Sliding window

## Tables

- Preserve headers
- Keep context
- Include table titles

## Code

- Function-level chunks
- Class-level chunks
- Keep imports

## Specialized

- Entity index
- Time-series index
- KPI index

# Tool-Based Approach: Interfaces

**Key concept:** Define specialized search tools with clear interfaces

## Tool Design Principles

- Define clear interfaces for each specialized retrieval method
  - Consistent input/output formats
  - Well-defined parameters and options
  - Clear documentation and examples

### Example Interface:

```
table_search(query, date_range,  
metrics=["revenue", "margin"])
```



### Specialized Retrieval: Beyond Basic Search

- **Implement hybrid search** combining lexical and semantic approaches
  - Lexical: Exact matches, codes, IDs, specific terms
  - Semantic: Concepts, themes, intentions, synonyms
- **Extract rich metadata** during document ingestion
- **Create specialized indices** for different content types
  - Text, tables, code, images each need different handling
- **Design tool-based interfaces** for specialized retrieval methods

**Remember:** Different content types  
require different retrieval strategies

# Tool-Based Approach: Routing

## Query Router

- Create a router that selects appropriate tool(s) for each query
  - Use LLM to classify and route queries
  - Consider confidence scores for tool selection
  - Fall back to general search when uncertain

## Execution Strategy

- Consider parallel execution for better performance
  - Run multiple tools simultaneously when appropriate
  - Merge results with intelligent ranking
  - Balance latency vs. thoroughness

# Implementing Query Routing

## Start Simple

- Begin with 3-5 core tools
- Focus on common query types
- Add more as patterns emerge

## Consistency Matters

- Standardize tool interfaces
- Common parameter formats
- Predictable outputs

**Testing tip:** Create synthetic queries  
for each tool to validate router accuracy

# Measuring Router Performance

- Test routing accuracy with synthetic queries
  - Create test cases for each tool
  - Measure routing precision and recall
  - Identify confusion patterns between tools
- Measure tool recall: Is the right tool being selected?
  - Track correct tool selection rate
  - Monitor unnecessary tool calls
  - Improve router prompts based on errors

## table\_search

- Finds financial tables
- Preserves row/column context
- Extracts precise metrics
- Example: "Q2 operating margin"

## text\_search

- Retrieves narratives
- MD&A, risk factors
- Semantic paragraph search
- Example: "Revenue growth factors"

## Financial Domain Tools (2/2)

### `entity_lookup`

- Company profiles
- Key metrics
- Industry data
- Example: "Company X's position"

### `time_series`

- Historical metrics
- Trend analysis
- Period comparisons
- Example: "Revenue growth over 8 quarters"

**Start with tools that address  
your most common queries**

# Case Study: Construction Blueprints & Visual Summaries

## The Challenge

- Construction company needed AI to answer blueprint questions
- Simple image captioning gave generic descriptions
- Only 27% recall on specific blueprint questions
- Multimodal retrieval was ineffective

## The Solution

- Prompted for detailed descriptions
- "Count floors, label mechanical rooms, highlight windows"
- Added specialized bounding-box extraction
- Merged text + blueprint data

**Key Takeaway:** Specific extraction prompts improved recall from 27% to 75-85% in just days

### Query Routing: Directing Queries to the Right Tools

- **Implement a query router** to direct questions to specialized tools
  - Use chain-of-thought prompting for better classification
  - Consider multi-tool routing for complex queries
- **Design clear tool interfaces** with consistent parameters
- **Balance precision and recall** in routing decisions
- **Implement fallback mechanisms** for uncertain classifications
- **Track routing accuracy** as a key performance metric

**Remember:** The right tool for the right job dramatically improves results



### Embeddings & Reranking: Optimizing Relevance

- **Choose embedding models** based on domain and query types
  - General models for broad topics
  - Domain-specific models for specialized fields
- **Implement reranking** to improve precision
  - Cross-encoders for higher accuracy
  - LLM-based reranking for complex relevance judgments
- **Fine-tune embeddings** with domain-specific feedback
- **Optimize chunk size** for your specific content and queries

**Remember:** Better embeddings and reranking can dramatically improve relevance

## Simple Binary Feedback with Categories

### Feedback UI

- Thumbs up/down buttons
- Simple and prominent
- Quick to complete

### Reason Categories

- Irrelevant results
- Incomplete answer
- Incorrect information
- Outdated content

**Key insight:** Even simple feedback is better than none

# User Feedback: Citation & Implicit Signals

## Citation Validation

- "Were these sources helpful?"
- Mark irrelevant citations
- Track valuable sources
- Identify missing citations

## Implicit Signals

- Time spent reviewing results
- Follow-up questions
- Copied/saved content
- Repeated queries

**Pro tip:** Combine explicit and implicit signals for a complete picture

## UI Elements

- Large, obvious buttons
- Modal dialogs for higher engagement
- 1-2 clicks for initial feedback
- Optional detailed feedback

## Feedback Categories

- Retrieval vs. generation issues
- Missing vs. incorrect info
- Technical errors vs. content gaps
- UI/UX problems

**Did you know?** Modal dialogs get 4-5x more feedback than subtle buttons

## Store Complete Context

- Original query and retrieved documents
- Generated response and citations
- User feedback and follow-up actions
- System metadata (latency, model version, etc.)

### Example schema:

```
{query, results, response,  
  feedback, metadata}
```

## Weekly Review

- Analyze feedback patterns
- Track trends over time
- Identify common failures
- Share insights with team

## Prioritization

- Focus on high-volume issues
- Balance quick wins vs. structural fixes
- Create targeted test cases
- Track impact of fixes

# Continuous Improvement Process

## Creating Test Cases

- Convert real failures into test examples
- Expand test coverage based on user behavior
- Validate fixes against expanded test set

## Model Refinement

- Use feedback to create training pairs
- Fine-tune embeddings or rerankers
- Improve router accuracy with real examples

**Key milestone:** When feedback  
drives automatic system improvements

# Case Study: Rejected URLs in Sales Follow-Up

## The Challenge

- System wrote post-call follow-up emails with links
- AI kept hallucinating or mistyping certain URLs
- 4% error rate on links in emails
- Damaged credibility with customers

## The Solution

- Introduced URL validator
- Rejected links to non-existent pages or unknown domains
- Asked model to remove/fix invalid links
- Later fine-tuned model to avoid invalid links

**Key Takeaway:** Simple post-processing + feedback loop reduced errors to nearly 0%



# Case Study: Changing Copy to Collect More Feedback

## The Challenge

- Team wanted more thumbs up/down data
- Tiny button labeled "How did we do?" hidden at top
- Very few users provided ratings
- Insufficient data for measuring correctness

## The Solution

- Changed copy to "Did we answer your question?"
- Made feedback UI larger and more central
- Requested brief reason for thumbs down
- Categorized feedback for targeted improvements

**Key Takeaway:** UI changes increased feedback volume 4-5x, providing crucial data for improvements

# Streaming Responses: Improving Perceived Performance

## Benefits of Streaming

- Reduces perceived latency by 11%
- Users tolerate longer wait times
- Allows immediate reading while generation continues
- Provides visual feedback on progress

## Implementation Approaches

- Stream tokens as they're generated
- Show interstitials during processing steps
- Render skeleton screens during loading
- Display tool execution in real-time

**Did you know?** Users will wait up to 8 seconds longer when given visual feedback

# Streaming Implementation: Interstitials & UI

- **Interstitial messages** during processing steps:
  - "Thinking..." → "Searching documents..." → "Reading results..." → "Generating response..."
- **Structured streaming** for complex responses:
  - Stream partial JSON objects with content, citations, follow-ups
  - Parse and render components as they arrive
- **Tool execution visualization:**
  - Show function calls and parameters in real-time
  - Allow user edits to parameters before execution

**Key insight:** Streaming isn't just about tokens—it's about communicating process

# Rejecting Work: Setting Expectations

## The Problem

- Not all queries can be answered well
- 90% success rate means 10% failures
- Attempting all queries reduces trust
- Users blame system for poor answers

## The Solution

- Identify query types with low success
- Gracefully reject answering these
- Set clear expectations with users
- Collect feedback on rejected queries

**Example:** "I can't confidently answer this question with the information available. Would you like me to try anyway with the understanding that the answer may be incomplete?"

# Showcasing Capabilities

## ■ Highlight what your system does well:

- Suggest example queries users can try
- Demonstrate different capabilities through UI
- Show different content types you can handle

## ■ Implementation approaches:

- Categorized example queries on landing page
- "Focus" buttons to expand capability options
- Follow-up suggestions after each response
- Special UI components for different content types

**Key insight:** Users will use the capabilities you highlight and ignore those you don't

# Chain of Thought: Improving Reasoning

## Benefits

- 10% performance improvement
- Makes complex reasoning possible
- Provides transparency into model thinking
- Enables better error analysis

## Implementation

- Structure as XML components
- Stream thinking process separately
- Allow users to expand/collapse
- Collect feedback on reasoning errors

**Key insight:** Chain of thought can be the difference between usable and unusable

# Monologuing: Managing Complex Contexts

- **Reiterate relevant information** before generating responses:

- Restate key instructions from the prompt
- Summarize relevant parts of the context
- Co-locate related information for better attention

- **Multi-stage monologuing** for complex tasks:

- Identify relevant variables first
- Extract relevant sections from documents
- Connect information across sources
- Reason about options before generating final response

**Case study:** SaaS pricing quotes improved dramatically with 4-stage monologuing

### Product Considerations: UX, Feedback & Prompting

- **Design prominent feedback mechanisms**
  - Make feedback UI large and obvious
  - Use modal dialogs for higher engagement
- **Implement streaming responses**
  - Reduce perceived latency with visual feedback
  - Use interstitials to communicate processing steps
- **Know when to reject work**
  - Set clear expectations for low-confidence answers
  - Showcase capabilities you excel at
- **Leverage chain of thought & monologuing**
  - Improve reasoning with structured thinking
  - Reiterate key information for better context handling



# Implementation Roadmap: Initial Phase

## Weeks 1-2: Evaluation Foundation

### Data Generation

- Create synthetic test data
- Cover diverse query types
- Include domain-specific examples

### Metrics Setup

- Implement evaluation pipeline
- Establish baseline metrics
- Build performance dashboard

## Weeks 3-4: Segmentation & Analysis

### Query Classification

- Define query categories
- Classify existing queries
- Create segment-specific test sets

### Performance Analysis

- Analyze metrics by segment
- Identify performance gaps
- Prioritize improvement areas

## Weeks 5-6: Specialized Retrieval

### Index Development

- Build domain-specific indices
- Optimize chunking strategies
- Implement metadata extraction

### Routing System

- Create basic query router
- Design tool interfaces
- Test with synthetic queries

# Implementation Roadmap: Refinement Phase

## Weeks 7-8: Feedback & Refinement

### Feedback Collection

- Add user feedback mechanisms
- Collect structured feedback
- Analyze feedback patterns

### Continuous Improvement

- Implement fine-tuning process
- Create improvement cycle
- Schedule regular reviews

**Key milestone:** Complete end-to-end system with feedback loop

# Implementation Roadmap: Key Insights

## Systematic Implementation: Building the Flywheel

- **Start with evaluation infrastructure** (Weeks 1-2)
  - Create synthetic test data
  - Establish baseline metrics
- **Analyze and segment queries** (Weeks 3-4)
  - Classify query types
  - Identify performance gaps
- **Build specialized retrieval** (Weeks 5-6)
  - Develop domain-specific indices
  - Implement query routing
- **Implement UX improvements & feedback loops** (Weeks 7-8)
  - Design prominent feedback mechanisms
  - Implement streaming responses
  - Add chain of thought reasoning
  - Set clear expectations with users

# RAG Implementation: The Complete Picture

## From Evaluation to Continuous Improvement

### Technical Foundation

- Evaluation metrics & test data
- Query segmentation & analysis
- Specialized retrieval methods
- Hybrid search & metadata

### User Experience

- Feedback collection mechanisms
- Streaming & interstitials
- Chain of thought reasoning
- Setting clear expectations

### Key Success Factors:

- Start with evaluation infrastructure before building features
- Segment queries to identify specific improvement areas
- Build specialized tools for different content types
- Create feedback loops to drive continuous improvement

# Key Resources & References

## Further Learning & Implementation Resources

### Technical Resources

- LangChain - Framework for RAG applications
- LlamaIndex - Data framework for LLM applications
- RAGAS - Evaluation framework for RAG systems
- LangChain Hub - Community prompts and chains

### Research & Best Practices

- RAG Survey Paper - Comprehensive overview
- Self-RAG - Self-reflective retrieval
- Chain-of-Thought - Reasoning techniques
- Instructor - Structured outputs

**Contact:** [jxn1.co](https://jxn1.co) — [@jxn1co](https://twitter.com/jxn1co) — [github.com/jxn1](https://github.com/jxn1)

## Conclusion & Next Steps

---



# Key Takeaways

## ■ Measurement First, Always

- Start with clear metrics before making changes
- Focus on retrieval recall as a leading indicator
- Build fast evaluation cycles with synthetic data

## ■ Systematic Over Ad-hoc

- Segment queries to focus improvements
- Build specialized tools for different content types
- Create feedback loops for continuous improvement

## ■ Domain Expertise Matters

- Generic embeddings aren't enough for specialized domains
- Fine-tune with domain-specific feedback
- Balance technical implementation with domain knowledge

# Common Pitfalls to Avoid

## ■ Intervention Bias

- Adding more prompt engineering without measurement
- Accumulating technical debt through unproven additions
- Focusing on generation quality before fixing retrieval

## ■ Absence Blindness

- Missing retrieval problems because they're less visible
- Not measuring what you can't directly observe
- Assuming generation issues when retrieval is the problem

## ■ Premature Optimization

- Fine-tuning models before collecting enough data
- Building complex architectures before proving value
- Optimizing for edge cases before handling common cases

# Resources & Tools

## ■ Evaluation Frameworks

- RAGAS: Comprehensive RAG evaluation toolkit
- LangChain Evaluation: Built-in metrics for RAG systems
- TruLens: Feedback-driven evaluation for LLM applications

## ■ Retrieval & Embedding Tools

- Sentence Transformers: Fine-tunable embedding models
- Cohere Rerank: Powerful reranking capabilities
- Weaviate/Pinecone/Qdrant: Vector databases with hybrid search

## ■ Feedback Collection

- Promptfoo: Prompt testing and evaluation
- Argilla: Data collection and annotation platform
- Weights & Biases: Experiment tracking and visualization

# Questions to Guide Your Implementation

- 1 What are your most critical query types and how will you measure success for each?
  - 2 How will you generate synthetic test data that matches your domain?
  - 3 What specialized retrieval methods do you need for your content types?
  - 4 How will you collect and incorporate user feedback?
  - 5 What's your plan for continuous improvement over time?
- Remember: The goal is to build a system that gets better with every interaction
  - Start simple, measure thoroughly, and improve systematically
  - Focus on the highest impact areas first based on data, not intuition

# Contact & Follow-up

## ■ Office Hours

- Available for project-specific questions
- Schedule via email: [jason@jxnl.co](mailto:jason@jxnl.co)

## ■ Additional Resources

- Course materials and code examples: [github.com/jxnl/mit-rag](https://github.com/jxnl/mit-rag)
- Reference implementations for different domains
- Recommended reading and tutorials: [jxnl.co/writing](https://jxnl.co/writing)

## ■ Community

- Join the discussion group for ongoing support
- Share your progress and learnings
- Connect with others working on similar challenges

Thank you! Questions?